

# FinalRMDScript\_050618\_ForPDF

Christian Conroy

May 6, 2018

## The Data

For this project, we use data provided by Yelp for the Yelp Dataset challenge, which can be found in JSON format from this website:

<https://www.yelp.com/dataset/challenge>

Documentation describing the variables and information contained in each json file comprising the yelp challenge dataset can be found here:

<https://www.yelp.com/dataset/documentation/json>

We then merge it with Zillow data on rental values across the United States by postal code. The Zillow rental value dataset can be downloaded from this website:

<https://www.zillow.com/research/data/>

## Setting the Working Directory and loading in required packages

### Importing the data

First, we import the Yelp checkin data and flatten it.

```
# Stream in Checkin Data
yelp_checkin <-
as.data.frame(jsonlite::stream_in(file("dataset/checkin.json")), flatten =
TRUE)
# Flatten Checkin Data
renquote <- function(l) if (is.list(l)) lapply(l, renquote) else enquote(l)
yelp_checkin_flat <- as.data.frame(lapply(unlist(renquote(yelp_checkin)),
eval))
```

### Reshaping the data

We clean the time period variable names by collapsing to long and using string functions to isolate day of the week name in string.

```
# Convert from wide to Long
yelp_checkin_flat_long <- reshape(yelp_checkin_flat, varying =
list(names(yelp_checkin_flat[1:168])), times =
names(yelp_checkin_flat[1:168]), idvar = 'business_id', v.names = 'checkin' ,
direction = 'long')
```

```

# Eliminate punctuation and digits
yelp_checkin_flat_long$time <- str_replace(yelp_checkin_flat_long$time,
"time.", "")
yelp_checkin_flat_long$time <- gsub('[:digit:]]+', '',
yelp_checkin_flat_long$time)

# Isolate name of weekday
yelp_checkin_flat_long$time =
substr(yelp_checkin_flat_long$time,1,nchar(yelp_checkin_flat_long$time)-2)

```

Here, we will collapse the data, reshape it from long to wide, then merge them together.

```

# Aggregate checkin by business and time period to get checkin average and
totalby business for each day of the week.
yelp_checkin_collapse_mean <- as.data.frame(aggregate(checkin ~ business_id +
time, yelp_checkin_flat_long , mean))
yelp_checkin_collapse_sum <- as.data.frame(aggregate(checkin ~ business_id +
time, yelp_checkin_flat_long , sum))

# Convert from Long to wide
yelp_checkin_wide_mean <- spread(yelp_checkin_collapse_mean, key = time,
value = checkin)
yelp_checkin_wide_sum <- spread(yelp_checkin_collapse_sum, key = time, value
= checkin)

# Merge averages and totals
yelp_checkin_wide <- inner_join(yelp_checkin_wide_mean,
yelp_checkin_wide_sum, by='business_id', match='all')
colnames(yelp_checkin_wide) <- c("business_id", "Friday_ave", "Monday_ave",
"Saturday_ave", "Sunday_ave", "Thursday_ave", "Tuesday_ave", "Wednesday_ave",
"Friday_total", "Monday_total", "Saturday_total", "Sunday_total",
"Thursday_total", "Tuesday_total", "Wednesday_total")

```

**Loading in business dataset to merge with check in data and aggregate by zip code. Then we eliminate useless columns and merge.**

```

# Import business dataset
yelp_business <- fromJSON(sprintf("[%s]",
paste(readLines("dataset/business.json"), collapse=",")),
simplifyDataFrame=TRUE, flatten=TRUE)

# Merge checkin data with business data by business
checkinbiz <- inner_join(yelp_business, yelp_checkin_wide,
by=c('business_id'), match='all')

# Eliminate unneeded columns
checkinbiz <- checkinbiz[-c(2:6, 8:101)]

# Collapse checkin data by zipcode to get total and average checkins each

```

*weekday for each zipcode*

```
checkinzipmean <- as.data.frame(aggregate(. ~ postal_code, checkinbiz[2:9],
mean))
checkinzipsum <- as.data.frame(aggregate(. ~ postal_code, checkinbiz[c(2,
10:16)], sum))

checkinfull <- inner_join(checkinzipmean, checkinzipsum, by=c('postal_code'),
match='all')
```

## Importing the Yelp Review Data

Because of the large size of the yelp review JSON file, The Yelp Review dataset was collapsed to the zipcode level through merging, aggregation, and collapsing within the google cloud platform.

*# Import data*

```
yelp_review_long <- read.csv("yelp_longPC_updated2.csv", header = T,
na.strings=c("NA"))
```

*# Check missingness*

```
apply(yelp_review_long, function(x) sum(is.na(x)))
```

```
##           X           postal_code           YearMonth
##           0              0              0
## Number_of_businesses Number_of_reviews           starsav
##           0              0              0
##           starssd           usefulav           funnyav
##          168953              0              0
##           coolav           bizstars           bizstarssd
##           0              0           168953
##           bizrevcount           bizrevav           is_openave
##           0              0              0
```

*# Check how many unique zipcodes*

```
length(unique(yelp_review_long$postal_code))
```

```
## [1] 15890
```

*# Convert YearMonth from character to ym class*

```
yelp_review_long$YearMonth <- as.yearmon(yelp_review_long$YearMonth)
```

Notice that there are no missing postal code values with 15,980 unique postal codes.

## Zillow data

Here, we read in the Zillow data that has information on all rental values across the US and Canada. We will rename the RegionName variable to "postal\_code" then convert the data from wide to long in order to merge it with the Yelp dataset. We will change the "time" variable to a date class, and finally we cut the Zillow data to match the dates of the Yelp

data (while Zillow data goes back to the 1990s, Yelp business and review data only go back to 2010).

```
# Import data
zillow <- read_csv("zecon/Zip_Zri_AllHomesPlusMultifamily.csv", col_names =
TRUE)

# Reformat to prepare for merge with Yelp dataset.
names(zillow)[2] <- "postal_code"
zillow <- as.data.frame(zillow)
zillow_long <- reshape(zillow, varying = list(names(zillow[8:95])), times =
names(zillow[8:95]), idvar = 'postal_code', v.names = 'rentprice', direction
= 'long')

sapply(zillow_long, function(x) sum(is.na(x)))

##      RegionID postal_code      City      State      Metro      CountyName
##           0           0           0           0      113960           0
##      SizeRank      time      rentprice
##           0           0      28354

zillow_long$time <- as.Date(strptime(paste(1, zillow_long$time), "%d %Y-%m"))

zillow_long <- zillow_long[zillow_long$time >= "2010-11-01" &
zillow_long$time <= "2017-12-31",]

zillow_long$month <- match(months(zillow_long$time), month.name)
zillow_long$year <- format(zillow_long$time, format="%Y")
zillow_long$YearMonth <- as.yearmon(paste(zillow_long$year,
zillow_long$month), "%Y %m")
names(zillow_long)

## [1] "RegionID"      "postal_code" "City"          "State"         "Metro"
## [6] "CountyName"    "SizeRank"    "time"          "rentprice"     "month"
## [11] "year"          "YearMonth"
```

Notice that there are 28,354 missing values. We will later attempt to rectify this through imputation.

## Merging all datasets

Here, we create the official merged dataset from which we conduct the analysis portion

```
# Merge Checkin with Yelp
yelp_review_long <- left_join(yelp_review_long, checkinfull,
by=c('postal_code'), match='all')

# Merge Zillow with Yelp
Full_data_long <- inner_join(yelp_review_long, zillow_long,
```

```

by=c('postal_code', 'YearMonth'), match='all')
length(unique(Full_data_long$postal_code))

## [1] 564

sapply(Full_data_long, function(x) sum(is.na(x)))

##           X           postal_code           YearMonth
##           0              0              0
## Number_of_businesses Number_of_reviews           starsav
##           0              0              0
##           starssd           usefullav           funnyav
##           2977              0              0
##           coolav           bizstars           bizstarssd
##           0              0              2977
##           bizrevcount           bizrevav           is_openave
##           0              0              0
##           Friday_ave           Monday_ave           Saturday_ave
##           1337              1337              1337
##           Sunday_ave           Thursday_ave           Tuesday_ave
##           1337              1337              1337
##           Wednesday_ave           Friday_total           Monday_total
##           1337              1337              1337
##           Saturday_total           Sunday_total           Thursday_total
##           1337              1337              1337
##           Tuesday_total           Wednesday_total           RegionID
##           1337              1337              0
##           City           State           Metro
##           0              0              64
##           CountyName           SizeRank           time
##           0              0              0
##           rentprice           month           year
##           580              0              0

```

The final frame consists of 37492 observations and 39 variables.

## Explortatory Data Analysis

Here is where we conduct the official analysis portion of the project

```

# First, we subset business based on eventual merging with Zillow
Full <- unique(Full_data_long$postal_code)
business_zillow <- dplyr::filter(yelp_business, postal_code %in% Full)


# Assessing the count of states
Full_data_long$State <- Full_data_long$State %>% as.factor
Full_data_long$State %>% summary

##   AR   AZ   CA   FL   GA   IL   IN   KY   LA   MD   MI   MO
##  22 10234 242   4   17  925  15  44  394   6   3   8
##  MS   NC   NJ   NM   NV   NY   OH  OK   PA   SC  TX  VA

```

```
##      4  4088      7      3  4981      75  7301      3  6510  405      17      6
##      VT      WI
##      7  2171
```

```
ggplot(Full_data_long, aes(State)) + geom_bar() + labs(title= " Frequency of
States in Yelp Challenge Dataset")
```



The states most represented in the dataset are Arizona, Ohio, Pennsylvania, Nevada, North Carolina, Wisconsin, and Illinois.


```
Full_data_long$City <- Full_data_long$City %>% as.factor
Full_data_long$City %>% summary
```

```
##      Phoenix      Las Vegas      Charlotte      Pittsburgh
##      3918      3580      1972      1847
##      Cleveland      Mesa      Madison      Scottsdale
##      1232      1115      1021      774
##      Glendale      Henderson      North Las Vegas      Chandler
##      695      691      522      516
##      Gilbert      Surprise      Peoria      Tempe
##      516      388      344      344
##      Champaign      Fort Mill      Gastonia      Shaler
##      258      256      251      237
##      Euclid      Lorain      Avondale      Concord
##      206      192      172      172
##      Cuyahoga Falls      Goodyear      Parma      Urbana
##      172      172      172      172
```

##	O'Hara	Ross	Chagrin Falls	Penn Hills
##	170	168	166	166
##	Sun City	Strongsville	Kannapolis	Queen Creek
##	159	158	156	119
##	Grafton	Amherst	Anthem	Avon
##	100	86	86	86
##	Bedford	Bellevue	Belmont	Berea
##	86	86	86	86
##	Bethel Park	Boulder City	Brecksville	Broadview Heights
##	86	86	86	86
##	Brook Park	Carefree	Carnegie	Castle Shannon
##	86	86	86	86
##	Cleveland Heights	Cornelius	Davidson	Elyria
##	86	86	86	86
##	Fairlawn	Fairview Park	Fountain Hills	Garfield Heights
##	86	86	86	86
##	Harrisburg	Hudson	Huntersville	Indian Trail
##	86	86	86	86
##	Kent	Lakewood	Litchfield Park	Lyndhurst
##	86	86	86	86
##	Malvern	Matthews	Medina	Mentor
##	86	86	86	86
##	Middleton	Mint Hill	Monroeville	Moon
##	86	86	86	86
##	Mount Charleston	Munhall	New Iberia	North Huntingdon
##	86	86	86	86
##	North Olmsted	North Ridgeville	North Strabane	Northfield
##	86	86	86	86
##	Paradise Valley	Parma Heights	Pineville	Richfield
##	86	86	86	86
##	Richmond Heights	Rocky River	Savoy	Seven Hills
##	86	86	86	86
##	Shaker Heights	Solon	South Euclid	Stallings
##	86	86	86	86
##	Stow	Streetsboro	Sun Prairie	(Other)
##	86	86	86	8891

```
temp <- row.names(as.data.frame(summary(Full_data_long$City, max=12))) #
create a df or something else with the summary output.
Full_data_long$City <- as.character(Full_data_long$City) # IMPORTANT! Here
was the problem: turn into character values
Full_data_long$top <- ifelse(
  Full_data_long$City %in% temp, ## condition: match aDDs$answer with
row.names in summary df
  Full_data_long$City, ## then it should be named as aDDs$answer
  "Other" ## else it should be named "Other"
)
Full_data_long$top <- as.factor(Full_data_long$top) # factorize the output
again
ggplot(Full_data_long[Full_data_long$top!="Other",], aes(x=factor(top,
```

```
levels=names(sort(table(top),increasing=TRUE)))) + geom_bar() +
labs(title="Frequency of Top Cities in Yelp Challenge Dataset") +
xlab("City") + ylab("Count")
```




This assesses the count of cities in the dataset. Top cities include Phoenix, Las Vegas, Charlotte, Pittsburgh, Cleveland.

```
# Reformat data to suit plot
catplot <- business_zillow%>%select(-starts_with("hours"), -
starts_with("attribute")) %>% unnest(categories) %>%
  select(name,
categories)%>%group_by(categories)%>%summarise(n=n())%>%arrange(desc(n))%>%he
ad(20)
catplot <- as.data.frame(catplot)

ggplot(data=catplot, aes(x=categories, y=n)) +
  geom_bar(stat="identity") + theme(axis.text.x = element_text(angle = 90,
hjust = 1)) + labs(title="Frequency of Categories in Dataset") +
xlab("Category") + ylab("Count")
```





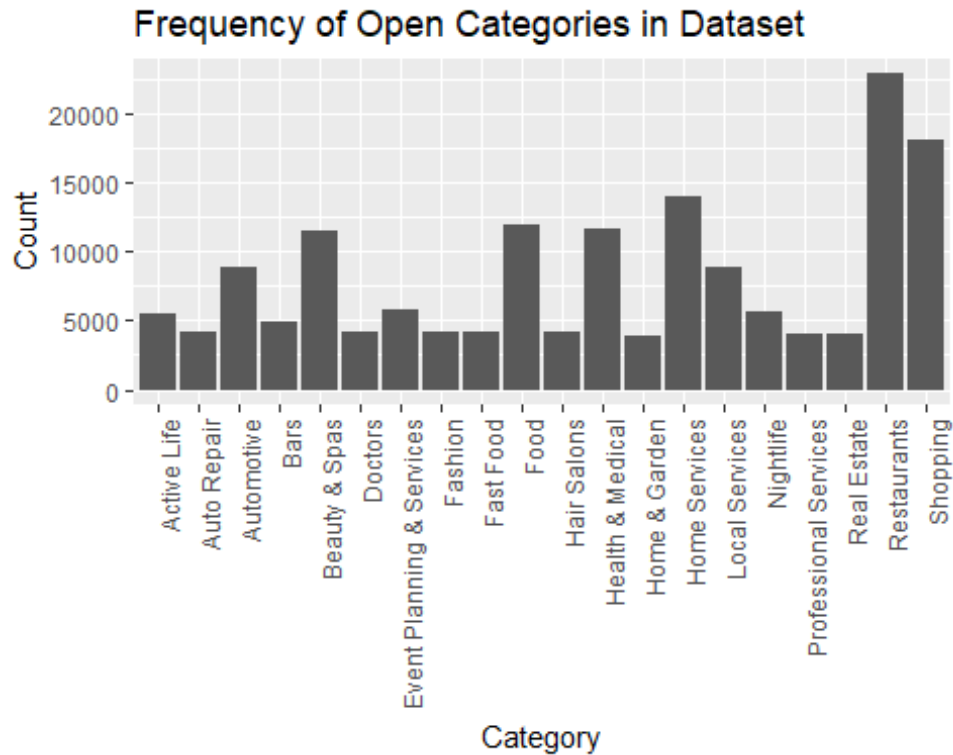
This counts the “categories” included in the dataset. These categories are tags users might use to describe various businesses. We can see that the top category is “Restuarants” followed by “Shopping,” “Food,” “Home Services,” and “Beauty and Spa.”

The Yelp dataset includes information on businesses that may have been open but are currently closed. The previous analyses included all of them, but here we assess the counts of categories only for businesses that are open.

```
catplot_open <- business_zillow %>%
  select(-starts_with("hours"), -starts_with("attribute")) %>%
  filter(is_open==1) %>%
  unnest(categories) %>%
  select(name, categories) %>%
  group_by(categories) %>%
  summarise(n=n()) %>%
  arrange(desc(n)) %>%
  head(20)

catplot_open <- as.data.frame(catplot_open )

ggplot(data=catplot_open , aes(x=categories, y=n)) +
  geom_bar(stat="identity") + theme(axis.text.x = element_text(angle = 90,
hjust = 1)) + labs(title="Frequency of Open Categories in Dataset") +
  xlab("Category") + ylab("Count")
```



Top categories include: Restaurants, Shopping, Home Services, Food, Health & Medical. Interestingly, it seems like “Food” and “Beauty & Spas” might be closing at higher rate than other categories.

Now, we compare the rent prices across cities. We use cities as categorizing variable because there are far too many zip codes in the dataset.

```
# Reimport data for plotting
zillow <- read_csv("zecon/Zip_Zri_AllHomesPlusMultifamily.csv", col_names =
TRUE)
# Eliminate unneeded columns
zillow <- zillow[-c(94:95)]


# Convert from long to wide by city
zillow_collapse_wide <- zillow %>%
  group_by(City) %>%
  summarize_all(funs(mean))
names(zillow_long)

## [1] "RegionID"      "postal_code" "City"          "State"         "Metro"
## [6] "CountyName"    "SizeRank"    "time"          "rentprice"     "month"
## [11] "year"          "YearMonth"
```

```
# Collapse to long by State and time period.
zillow_collapse_long <- zillow_long %>%
  group_by(State, time) %>%
  summarize(rentprice = mean(rentprice))
```

Finally, we compose a graph illustrating rents over time

```
ggplot(zillow_collapse_long, aes(x = time, y=rentprice, group = State, colour = State)) + geom_line() + scale_colour_discrete(guide = 'none') + scale_x_date(expand=c(0.1, 0)) + geom_dl(aes(label = State), method = list(dl.trans(x = x + .2), "last.points")) + geom_dl(aes(label = State), method = list(dl.trans(x = x - .2), "first.points")) + labs(title = "Rent Prices over Time by State", xlab = "Rent Price (USD)")
```




Few states in the dataset appear to have experienced overall declines in rent over the periods in question. California, Oregon, Colorado, Massachussets, and Washington appear to have experienced significant increases over the time period. DC, California, New Jersey, Hawaii, and Massachussets are consistently plagued by high rents prices.

## Comparing the average star values across cities. First, we must collapse by state.

```
# First, we collapse by state
FDL_collapse_long <- Full_data_long %>%
  group_by(State, time) %>%
  summarize(starsav = mean(starsav))

ggplot(FDL_collapse_long, aes(x = time, y=starsav, group = State, colour = State)) + geom_line() + scale_colour_discrete(guide = 'none') + scale_x_date(expand=c(0.1, 0)) + geom_dl(aes(label = State), method = list(dl.trans(x = x + .2), "last.points")) + geom_dl(aes(label = State), method = list(dl.trans(x = x - .2),
```

```
"first.points")) + labs(title = "Average Stars over Time by State", xlab =  
"Avg. Stars (1-5)")
```



As we can see, there don't appear to be any trends or patterns. Seems like a very random relationship.

## Modeling and Panel Data Regression Analysis

Now comes the fun part. We split the data into train and test sets where our test set comprises the last year (12 months) of our data set. We run a Hausman test to determine whether we should run a fixed effects (“within”) or a random effects model. Then we develop a Panel Linear Regression model to predict housing prices.

[illegible]

```

                                index = my.index)
print(my.hausman.test.train)

##
## Hausman Test
##
## data: my.formula
## chisq = 1.2796, df = 5, p-value = 0.937
## alternative hypothesis: one model is inconsistent

```

The high p-value of 0.937 indicates that we should use a random effects model instead of a fixed effects model.

Now, we build our random effects model on the training dataset and predict on the test set. We calculate the Mean Average Percent Error (MAPE) to see how accurately our model uses training values to predict rental prices in the test set.

```

my.pdm.train <- plm(data = Full_data_long_train,
                    formula = my.formula,
                    model = 'random',
                    index = my.index)
summary(my.pdm.train)

## Oneway (individual) effect Random Effect Model
## (Swamy-Arora's transformation)
##
## Call:
## plm(formula = my.formula, data = Full_data_long_train, model = "random",
##      index = my.index)
##
## Unbalanced Panel: n = 560, T = 1-74, N = 31397
##
## Effects:
##              var   std.dev share
## idiosyncratic 6321.76    79.51 0.023
## individual    270893.25   520.47 0.977
## theta:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.8490 0.9817 0.9822 0.9806 0.9822 0.9822
##
## Residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -2126.91  -40.27   -6.88   -0.36   37.76  2068.60
##
## Coefficients:
##              Estimate Std. Error t-value Pr(>|t|)
## (Intercept) 1238.62701   22.65028  54.6848 < 2.2e-16 ***
## starsav      -2.85126    0.83058  -3.4329 0.000598 ***
## is_openave    131.20358    3.13631  41.8337 < 2.2e-16 ***
## funnyav      -10.17243    0.86013 -11.8266 < 2.2e-16 ***
## coolav        11.10536    0.96571  11.4997 < 2.2e-16 ***

```

```
## usefulav      -11.10133      0.42243 -26.2800 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Total Sum of Squares:    226790000
## Residual Sum of Squares: 203110000
## R-Squared:      0.10443
## Adj. R-Squared: 0.10429
## F-statistic: 731.838 on 5 and 31391 DF, p-value: < 2.22e-16


Full_data_long_test$pred.plm.test <- predict(my.pdm.train,
Full_data_long_test, type='response')

plmmape <-
100*mean(abs(Full_data_long_test$pred.plm.test/Full_data_long_test$rentprice-
1), na.rm = T)
print(plmmape)

## [1] 21.39346
```

MAPE is only 21.39% right now. We will continue working on the model to get this error lower.

```
ggplot(Full_data_long_test, aes(x=rentprice, y=pred.plm.test)) +geom_point()
+ labs(title="Predicted vs. Actual Real Estate Prices") + xlab("Actual") +
ylab("Predicted")
```



The plot above shows a significant discrepancy existing between actual and predicted rent prices.

## Generating the Lag Model

To fine-tune the model, we decide to lag the dependent variable to consider the possibility that last month's rent could be the best predictor of this month's rent price. We follow a similar process to the one above for training, testing, and predicting.

```
my.lag.formula <- rentprice ~ lag(rentprice, 1) + starsav + is_openave +  
funnyav + coolav + usefultav + Number_of_reviews
```

```
# Conduct Hausman Test
```

```
my.hausman.test.train.lag <- phtest(x = my.lag.formula,  
                                   data = Full_data_long_train,  
                                   model = c('within', 'random'),  
                                   index = my.index)
```

```
# print result
```

```
print(my.hausman.test.train.lag)
```

```
##  
## Hausman Test  
##  
## data: my.lag.formula  
## chisq = 289.55, df = 7, p-value < 2.2e-16  
## alternative hypothesis: one model is inconsistent
```

Then, we build the model on the training set and predict on the test set in order to calculate the MAPE.

```
my.pdm.train.lag <- plm(data = Full_data_long_train,  
                       formula = my.lag.formula,  
                       model = 'random',  
                       index = my.index)
```

```
summary(my.pdm.train.lag)
```

```
## Oneway (individual) effect Random Effect Model  
## (Swamy-Arora's transformation)  
##  
## Call:  
## plm(formula = my.lag.formula, data = Full_data_long_train, model =  
## "random",  
## index = my.index)  
##  
## Unbalanced Panel: n = 519, T = 1-73, N = 29252  
##  
## Effects:  
## var std.dev share  
## idiosyncratic 161.483 12.708 0.993  
## individual 1.107 1.052 0.007  
## theta:  
## Min. 1st Qu. Median Mean 3rd Qu. Max.
```

```
## 0.00341 0.17795 0.18361 0.17177 0.18361 0.18361
##
## Residuals:
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -142.410   -4.848     0.001     0.009     5.213    132.569
##
## Coefficients:
##              Estimate Std. Error t-value Pr(>|t|)
## (Intercept)  -0.81953342  0.73031005  -1.1222  0.261799
## lag(rentprice, 1)  1.00093673  0.00017581 5693.3592 < 2.2e-16 ***
## starsav       -0.25238363  0.14094489  -1.7907  0.073359 .
## is_openave     3.34976459  0.48784817   6.8664 6.715e-12 ***
## funnyav       -0.49065915  0.15640230  -3.1372  0.001708 **
## coolav        -0.01465219  0.16578108  -0.0884  0.929573
## usefulav      -0.19242788  0.07260164  -2.6505  0.008043 **
## Number_of_reviews 0.00124545  0.00029642   4.2016 2.659e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Total Sum of Squares:    5351700000
## Residual Sum of Squares: 4771300
## R-Squared:    0.99911
## Adj. R-Squared: 0.99911
## F-statistic: 4681730 on 7 and 29244 DF, p-value: < 2.22e-16

# Predict
Full_data_long_test$pred.plm.test.lag <- predict(my.pdm.train.lag,
Full_data_long_test, type='response')

# MAPE
plmmape.lag <-
100*mean(abs(Full_data_long_test$pred.plm.test.lag/Full_data_long_test$rentpr
ice-1), na.rm = T)
print(plmmape.lag)

## [1] 0.2112457
```

Now we get a MAPE of 0.211, far lower than the non-lagged model. This supports the hypothesis that last month's rent could be the best predictor of this month's rent price.

```
ggplot(Full_data_long_test, aes(x=rentprice, y=pred.plm.test.lag))
+geom_point() + labs(title="Predicted vs. Actual Real Estate Prices") +
xlab("Actual") + ylab("Predicted")
```





## Multiple Imputation for Missing Values Using the Amelia Package

This process uses bootstrapping and Expectation-Maximization algorithm to impute the missing values in a data set. In our model, we will be able to throw in almost all of our independent variables.

*# Look at missingness to get a sense of what needs to be imputed.*

```
sapply(Full_data_long, function(x) sum(is.na(x)))
```

```
##           X           postal_code           YearMonth
##           0                0                0
## Number_of_businesses Number_of_reviews           starsav
##           0                0                0
##           starssd           usefulav           funnyav
##           2977                0                0
##           coolav           bizstars           bizstarssd
##           0                0                2977
##           bizrevcount           bizrevav           is_openave
##           0                0                0
##           Friday_ave           Monday_ave           Saturday_ave
##           1337           1337           1337
##           Sunday_ave           Thursday_ave           Tuesday_ave
##           1337           1337           1337
##           Wednesday_ave           Friday_total           Monday_total
##           1337           1337           1337
##           Saturday_total           Sunday_total           Thursday_total
```

```
##          1337          1337          1337
##      Tuesday_total      Wednesday_total      RegionID
##          1337          1337          0
##          City          State          Metro
##          0          0          64
##      CountyName          SizeRank          time
##          0          0          0
##      rentprice          month          year
##          580          0          0
##          top
##          0
```

```
Full_data_long <- Full_data_long[-c(40)]
Imputed_Full_data_long <- amelia(Full_data_long, ts= 'time', cs= 'postal_code',
p2s=0, intercs = FALSE, idvars=c('City', 'State', 'Metro', 'CountyName',
'year', 'month', 'YearMonth'))
```

```
write.amelia(obj=Imputed_Full_data_long, file.stem="imputedfull")
```

```
data1 <- read.csv("imputedfull1.csv")
data2 <- read.csv("imputedfull2.csv")
data3 <- read.csv("imputedfull3.csv")
data4 <- read.csv("imputedfull4.csv")
data5 <- read.csv("imputedfull5.csv")
```

```
data1 <- pdata.frame(data1, index = c("postal_code", "time"))
data2 <- pdata.frame(data2, index = c("postal_code", "time"))
data3 <- pdata.frame(data3, index = c("postal_code", "time"))
data4 <- pdata.frame(data4, index = c("postal_code", "time"))
data5 <- pdata.frame(data5, index = c("postal_code", "time"))
```

```
allimp <- imputationList(list(data1,data2,data3,data4,data5))
```

Now, we will create the train and tests set using the last 12 months (1 year) for the test set, but with imputed values from an Amelia imputation iteration.

```
data5$time <- as.Date(data5$time, "%Y-%m-%d")
data5_train <- data5[data5$time < "2017-01-01",]
data5_test <- data5[data5$time >= "2017-01-01",]
```

```
my.formula.impute.lag <- rentprice ~ lag(rentprice, 12) + starsav + starssd +
is_openave + funnyav + coolav + usefulav + Number_of_reviews +
Number_of_businesses + Friday_ave + Monday_ave + Saturday_ave + Sunday_ave +
Thursday_ave + Tuesday_ave + Wednesday_ave + Friday_total + Monday_total +
Saturday_total + Sunday_total + Thursday_total + Tuesday_total +
Wednesday_total
```

```
my.index <- c('postal_code','time')
```

```
# Conduct Hausman Test
```

```
my.hausman.test.train.impute.lag <- phtest(x = my.formula.impute.lag,
                                           data = data5_train,
                                           model = c('within', 'random'),
                                           index = my.index)

# print result
print(my.hausman.test.train.impute.lag)

##
## Hausman Test
##
## data: my.formula.impute.lag
## chisq = 29325, df = 23, p-value < 2.2e-16
## alternative hypothesis: one model is inconsistent
```

Build random effects model on train and predict on test.

```
my.pdm.train.impute.lag <- plm(data = data5_train,
                               formula = my.formula.impute.lag,
                               model = 'random',
                               index = my.index)
summary(my.pdm.train.impute.lag)

## Oneway (individual) effect Random Effect Model
## (Swamy-Arora's transformation)
##
## Call:
## plm(formula = my.formula.impute.lag, data = data5_train, model = "random",
##      index = my.index)
##
## Unbalanced Panel: n = 425, T = 1-62, N = 22754
##
## Effects:
##               var std.dev share
## idiosyncratic 4866.88   69.76 0.901
## individual    533.53   23.10 0.099
## theta:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.05068 0.64187 0.64187 0.62772 0.64187 0.64187
##
## Residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -2176.60  -31.74   -6.65    0.07   23.44  1316.70
##
## Coefficients:
##               Estimate Std. Error t-value Pr(>|t|)
## (Intercept)   135.3726454   13.4967249   10.0300 < 2.2e-16 ***
## lag(rentprice, 12)  0.8362233   0.0032052  260.8992 < 2.2e-16 ***
## starsav        4.3925615   2.0918703    2.0998 0.0357552 *
## starssd       12.4481104   2.9101549    4.2775 1.898e-05 ***
## is_openave     56.1246937   6.3694282    8.8116 < 2.2e-16 ***
## funnyav       -6.9686115   2.0246885   -3.4418 0.0005789 ***
```

```
## coolav          4.8148438    1.9313426    2.4930 0.0126738 *
## usefulav       -4.2287045    0.8987007   -4.7054 2.549e-06 ***
## Number_of_reviews -0.0168226    0.0103938   -1.6185 0.1055644
## Number_of_businesses 0.5307349    0.0437210   12.1391 < 2.2e-16 ***
## Friday_ave     -25.4383904    9.6657602   -2.6318 0.0084990 **
## Monday_ave     -13.9461762    7.8711624   -1.7718 0.0764401 .
## Saturday_ave    6.0077063    6.6304076    0.9061 0.3649010
## Sunday_ave     12.7093866    6.1359542    2.0713 0.0383423 *
## Thursday_ave   17.2457619    9.8452913    1.7517 0.0798430 .
## Tuesday_ave    -21.8732883   11.0744473   -1.9751 0.0482674 *
## Wednesday_ave  32.0943737   11.2596608    2.8504 0.0043706 ***
## Friday_total    0.0231701    0.0049521    4.6789 2.901e-06 ***
## Monday_total   -0.0218585    0.0043300   -5.0482 4.495e-07 ***
## Saturday_total -0.0242709    0.0040466   -5.9979 2.029e-09 ***
## Sunday_total    0.0221750    0.0037099    5.9772 2.303e-09 ***
## Thursday_total  0.0023054    0.0077629    0.2970 0.7664894
## Tuesday_total   0.0247671    0.0103663    2.3892 0.0168932 *
## Wednesday_total -0.0307317    0.0092509   -3.3220 0.0008950 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Total Sum of Squares:    1000700000
## Residual Sum of Squares: 215590000
## R-Squared:              0.78455
## Adj. R-Squared: 0.78434
## F-statistic: 3598.79 on 23 and 22730 DF, p-value: < 2.22e-16

data5_test$pred.plm.test.impute.lag <- predict(my.pdm.train.impute.lag,
data5_test, type='response')

plmmape_impute_lag <-
100*mean(abs(data5_test$pred.plm.test.impute.lag/data5_test$rentprice-1),
na.rm = T)
print(plmmape_impute_lag)

## [1] 5.102004
```

Imputation gives us 5.127882 (Might be different if we tried the other 4 imputed data sets)

```
ggplot(data5_test, aes(x=rentprice, y=pred.plm.test.impute.lag)) +
geom_point() + labs(title="Predicted vs. Actual Real Estate Prices") +
xlab("Actual") + ylab("Predicted")
```



Now, we conduct a reduced imputed model, which excludes checkin data

```
my.formula.impute.lag.Simple <- rentprice ~ lag(rentprice, 12) + starsav +
starsd + is_openave + funnyav + coolav + usefultav + Number_of_reviews +
Number_of_businesses
my.index <- c('postal_code', 'time')

my.hausman.test.train.impute.lag.Simple <- phptest(x =
my.formula.impute.lag.Simple,
data = data5_train,
model = c('within',
'random'),
index = my.index)

print(my.hausman.test.train.impute.lag.Simple)

##
## Hausman Test
##
## data: my.formula.impute.lag.Simple
## chisq = 30105, df = 9, p-value < 2.2e-16
## alternative hypothesis: one model is inconsistent
```

Build random effects model on train and predict on test

```
my.pdm.train.impute.lag.Simple <- plm(data = data5_train,
formula = my.formula.impute.lag.Simple,
```

```

                                model = 'random',
                                index = my.index)
summary(my.pdm.train.impute.lag.Simple)

## Oneway (individual) effect Random Effect Model
##   (Swamy-Arora's transformation)
##
## Call:
## plm(formula = my.formula.impute.lag.Simple, data = data5_train,
##      model = "random", index = my.index)
##
## Unbalanced Panel: n = 425, T = 1-62, N = 22754
##
## Effects:
##               var std.dev share
## idiosyncratic 4865.16   69.75 0.902
## individual    531.40   23.05 0.098
## theta:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.05051 0.64130 0.64130 0.62714 0.64130 0.64130
##
## Residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -2192.86  -32.02   -7.49    0.01   23.50  1328.23
##
## Coefficients:
##               Estimate Std. Error t-value Pr(>|t|)
## (Intercept)   114.1485383   12.1258561    9.4136 < 2.2e-16 ***
## lag(rentprice, 12)    0.8450445    0.0031533  267.9877 < 2.2e-16 ***
## starsav        6.2409575    2.0976999    2.9751 0.0029316 **
## starssd       16.2780616    2.9084258    5.5969 2.208e-08 ***
## is_openave     71.2143502    6.2535988   11.3877 < 2.2e-16 ***
## funnyav       -7.8383700    2.0344937   -3.8527 0.0001171 ***
## coolav         5.3923406    1.9406631    2.7786 0.0054637 **
## usefulav      -4.9796951    0.9009426   -5.5272 3.289e-08 ***
## Number_of_reviews -0.0266971    0.0083995   -3.1784 0.0014827 **
## Number_of_businesses 0.3370970    0.0359868    9.3672 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Total Sum of Squares:    1003100000
## Residual Sum of Squares: 218300000
## R-Squared:    0.78237
## Adj. R-Squared: 0.78228
## F-statistic: 9084.77 on 9 and 22744 DF, p-value: < 2.22e-16

# Predict
data5_test$pred.plm.test.impute.lag.Simple <-
predict(my.pdm.train.impute.lag.Simple, data5_test, type='response')

```

```
plmmape_impute_lag.Simple <-
100*mean(abs(data5_test$pred.plm.test.impute.lag.Simple/data5_test$rentprice-
1), na.rm = T)
print(plmmape_impute_lag.Simple)

## [1] 4.943544
```

## Imputation gives us 5.037

```
ggplot(data5_test, aes(x=rentprice, y=pred.plm.test.impute.lag.Simple))
+geom_point() + labs(title="Predicted vs. Actual Real Estate Prices") +
xlab("Actual") + ylab("Predicted")
```



## Final Model

The last thing we do is subset the Business dataset to include only businesses categorized as food or bars. We do this because we expect these businesses to have a stronger relationship to rent prices than others, such as Beauty & Spas.

There is not a significant change in the MAPE for the subset, with the non-imputed subset without a lagged dependent producing a MAPE of 18.69, the non-imputed subset with a lagged dependent producing a MAPE of .2142, and the imputed subset with a lag producing a MAPE of 3.46.