

Statistical Learning HW 3 - Neural Networks

Christian Conroy

March 1, 2019

Extra 28 (3 Points)

Below is the output from nnet after we fit a model. Let's assume we used a $\tanh()$ activation function throughout. Let $x_i, i = 1, 2, \dots$ be the input variables and let h_1, h_2, \dots be the output from the hidden layer.

```
a 2-2-1 network with 9 weights
options were - linear output units
b->h1 i1->h1 i2->h1
  1.2  4.2  -0.5
b->h2 i1->h2 i2->h2
-30    20   -40
b->o  h1->o  h2->o
  5    -8   1.5
```

Figure 1: Alt text

- (a) Draw a diagram of this neural network architecture. Label all the edges with the corresponding weights.
- (b) Provide an expression for the output value of the first hidden unit as a function of the values of the input features. This should have the form $h_1 = f(x_1, x_2, \dots)$ for a suitable explicit function f .

$$h_1 = f(1.2b + 4.2x_1 - 0.5x_2)$$

- (c) Provide an expression for the value at the output node as a function of the values at the hidden units. This should have the form $z = g(h_1, h_2, \dots)$ for a suitable explicit function g .

$$z = g(5b - 8h_1 + 1.5h_2)$$

- (d) Provide an expression for the value at the output node as a function of the input values. This should have the form $z = F(x_1, x_2, \dots)$ for a suitable explicit function F .

$$z = F(5b - 8(1.2b + 4.2x_1 - 0.5x_2) + 1.5(-30b + 20x_1 - 40x_2))$$

Extra 30 (3 points)

Suppose two different ANNs have been trained on a training set for a classification problem, and the responses, scaled to have values in $[0, 1]$, have been computed for all training instances for both networks. Assume that the responses for network 2 are related to those from network 1 by a monotone function, such as in the plot below. Explain carefully why the two ANNs have the same ROC curve.

Given that we are training both networks on the same train set, the sensitivity (True Positive) and 1 - specificity (False Positive) axes that comprise the ROC curve at different decision thresholds (the probability threshold considered to assign the positive class) will be the same for both networks. If the responses are related by a monotone function and the responses from the train set are scaled $[0,1]$, then we are essentially dealing with an ANN with a logistic activation function, or a logistic regression model.




Figure 2: Alt text

Extra 34 (3 points)

Make a dataframe with $k = 11$ columns and $N = 100$ observations, where all entries are independent standard normal random sample. Let z be the last column. Use `set.seed(20305)`.

```
set.seed(20305)
data <- data.frame(matrix(rnorm(100), nrow = 100, ncol = 11, byrow = TRUE))
colnames(data)[11] <- "z"
```

(a) Fit z to the other 10 columns using multiple regression. What is the sum of squares of the residuals?

```
reg1 <- lm(z ~ ., data = data)
summary(reg1)
```

```
##
## Call:
## lm(formula = z ~ ., data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.168 -0.546  0.039  0.549  2.645
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.1300     0.1084   1.20    0.23
## X1           -0.0284     0.1060  -0.27    0.79
## X2             0.0553     0.1060   0.52    0.60
## X3           -0.0834     0.1058  -0.79    0.43
## X4           -0.0886     0.1057  -0.84    0.40
## X5           -0.0663     0.1055  -0.63    0.53
## X6           -0.1215     0.1055  -1.15    0.25
## X7             0.0362     0.1057   0.34    0.73
## X8           -0.0708     0.1058  -0.67    0.50
## X9             0.0631     0.1060   0.60    0.55
## X10          -0.0598     0.1060  -0.56    0.57
##
## Residual standard error: 1.02 on 89 degrees of freedom
## Multiple R-squared:  0.0535, Adjusted R-squared:  -0.0528
## F-statistic: 0.503 on 10 and 89 DF,  p-value: 0.884
```

```
anova(reg1)
```

```
## Analysis of Variance Table
##
## Response: z
##          Df Sum Sq Mean Sq F value Pr(>F)
## X1         1    0.0   0.048    0.05  0.83
## X2         1    0.2   0.233    0.22  0.64
## X3         1    0.4   0.375    0.36  0.55
## X4         1    1.1   1.081    1.04  0.31
## X5         1    0.2   0.171    0.16  0.69
## X6         1    1.8   1.773    1.70  0.20
## X7         1    0.2   0.227    0.22  0.64
## X8         1    0.6   0.589    0.57  0.45
## X9         1    0.4   0.416    0.40  0.53
## X10        1    0.3   0.332    0.32  0.57
```

```
## Residuals 89    92.8    1.042
```

The sum of squared errors is 92.8.

- (b) Fit z to the other 10 columns, using a neural network with two hidden units and setting `maxit = 2000` and `decay = .01`. Does this model fit the data better? How do you know?

```
nn_fit <- nnet(z ~ ., data=data, maxit = 2000, decay = .01, size=2)
```

```
## # weights:  25
## initial  value 103.519040
## iter   10 value 98.170869
## iter   20 value 86.697084
## iter   30 value 83.853586
## iter   40 value 83.076781
## iter   50 value 83.032330
## iter   60 value 83.021933
## final   value 83.021895
## converged
```

This model fits the data better because we've added more hidden nodes, thereby making our model more flexible. The data converges (i.e. no longer fitting or doing gradient descent) at 83.02 (cross-entropy here is comparable to the SSR above), meaning we have a lower loss function than we did in the regression model (SSR of 92.8)

- (c) Redo this experiment with the same data and with 5 and 10 hidden units and explain what you see.

```
nn_fit <- nnet(z ~ ., data=data, maxit = 2000, decay = .01, size=5)
```

```
## # weights:  61
## initial  value 114.824303
## iter   10 value 99.389701
## iter   20 value 89.005620
## iter   30 value 84.418027
## iter   40 value 83.957663
## iter   50 value 79.743775
## iter   60 value 75.989300
## iter   70 value 75.301282
## iter   80 value 71.536431
## iter   90 value 64.858485
## iter  100 value 63.042621
## iter  110 value 62.147478
## iter  120 value 61.886012
## iter  130 value 61.764596
## iter  140 value 61.720205
## iter  150 value 61.713558
## iter  160 value 61.712330
## iter  170 value 61.712165
## final   value 61.712161
## converged
```

As expected, increasing the number of hidden nodes from 2 to 5 decreases the loss function even further to 63.97.

```
nn_fit <- nnet(z ~ ., data=data, maxit = 2000, decay = .01, size=10)
```

```
## # weights:  121
## initial  value 101.036142
## iter   10 value 96.993502
```

```
## iter 20 value 84.799456
## iter 30 value 83.746893
## iter 40 value 82.419975
## iter 50 value 81.107861
## iter 60 value 76.285676
## iter 70 value 75.588551
## iter 80 value 72.379225
## iter 90 value 63.615488
## iter 100 value 60.518175
## iter 110 value 60.013365
## iter 120 value 59.720134
## iter 130 value 59.416294
## iter 140 value 59.255925
## iter 150 value 59.211544
## iter 160 value 59.193762
## iter 170 value 59.116708
## iter 180 value 59.040275
## iter 190 value 59.016411
## iter 200 value 58.988584
## iter 210 value 58.980605
## iter 220 value 58.979042
## iter 230 value 58.978856
## iter 240 value 58.978837
## iter 240 value 58.978837
## iter 240 value 58.978837
## final value 58.978837
## converged
```

As expected, increasing the number of hidden nodes even further from 5 to 10 decreases the loss function even further to 58.35.

Extra 33 (5 points)

We'll use the MNIST image classification data, available as `mnist_all.RData` that were used in class during the last two weeks. We want to distinguish between 4 and 7. Extract the relevant training data and place them in a data frame.

```
data <- load('mnist_all.RData')
train <- data.frame(train$n, train$x, train$y)
train <- train[train$train.y == 4 | train$train.y == 7,]
test <- data.frame(test$n, test$x, test$y)
test <- test[test$test.y == 4 | test$test.y == 7,]
```

- (a) Pick two features (variables) that have large variances and low correlation. Fit a logistic regression model with these two features. Evaluate the model with the AUC score.

```
lvlc <- sort(sapply(train, function(x) sd(x)), decreasing = TRUE)[1:20]
lvlc <- train[,names(lvlc)]
cor(lvlc)
```


```
# X346 and x602 have high variance and low correlation.
```

```
spec <- as.formula("factor(train.y) ~ X346 + X602")
```

```
fit <- glm(formula = spec, data = train, family=binomial)

train$pred <- predict(fit, type = 'response')

r <- roc(train$train.y, train$pred)
plot(r)
```



```
auc(train$train.y, train$pred)
```

```
## Area under the curve: 0.772
```


The AUC is 0.772, meaning that it performs alright but not great.

- (b) Create a neural net with one unit in the hidden layer. Train the neural net with the same two features as the previous part and evaluate the model with AUC. Compare to the results from (a) and explain.

```
nn_fit <- nnet(formula = spec, data=train, maxit = 2000, decay = .01, size=1)
```

```
## # weights: 5
## initial value 8904.847464
## iter 10 value 6945.201478
## iter 20 value 6857.829755
## iter 30 value 6846.002394
## iter 40 value 6838.127984
## iter 50 value 6837.903772
## final value 6837.454464
## converged
```

```
train$pred_nn <- predict(nn_fit, type = "r")
r <- roc(train$train.y, train$pred_nn)
plot(r)
```



```
auc(train$train.y, train$pred_nn)
```

```
## Area under the curve: 0.745
```

The neural network model with one hidden layer actually does worse than the logistic model.

- (c) With the same two features, train three different neural nets, each time using more units in the hidden layer. How do the results improve, using the AUC?

```
nn_fit2 <- nnet(formula = spec, data=train, maxit = 2000, decay = .01, size=2)
```

```
## # weights: 9
## initial value 8607.652293
## iter 10 value 6787.220271
## iter 20 value 6744.200230
## iter 30 value 6659.992150
## iter 40 value 6652.132648
## iter 50 value 6651.295953
## iter 60 value 6651.034465
## iter 70 value 6650.984180
## final value 6650.982669
## converged
```

```
train$pred_nn2 <- predict(nn_fit2, type = "r")
```

```
auc(train$train.y, train$pred_nn2)
```

```
## Area under the curve: 0.782
```

```
nn_fit5 <- nnet(formula = spec, data=train, maxit = 2000, decay = .01, size=5)
```

```
## # weights: 21
```

```
## initial value 8551.820658
```

```
## iter 10 value 6855.649044
```

```
## iter 20 value 6693.498410
```

```
## iter 30 value 6641.121155
```

```
## iter 40 value 6634.420139
```

```
## iter 50 value 6632.113541
```

```
## iter 60 value 6622.491940
```

```
## iter 70 value 6611.453552
```

```
## iter 80 value 6608.768684
```

```
## iter 90 value 6607.941568
```

```
## final value 6607.793969
```

```
## converged
```

```
train$pred_nn5 <- predict(nn_fit5, type = "r")
```

```
auc(train$train.y, train$pred_nn5)
```

```
## Area under the curve: 0.788
```

```
nn_fit10 <- nnet(formula = spec, data=train, maxit = 2000, decay = .01, size=10)
```

```
## # weights: 41
```

```
## initial value 10175.176131
```

```
## iter 10 value 6724.726066
```

```
## iter 20 value 6646.309963
```

```
## iter 30 value 6631.225676
```

```
## iter 40 value 6617.736399
```

```
## iter 50 value 6611.523539
```

```
## iter 60 value 6609.441181
```

```
## iter 70 value 6605.720802
```

```
## iter 80 value 6605.308005
```

```
## iter 90 value 6604.615895
```

```
## iter 100 value 6603.922831
```

```
## iter 110 value 6603.682384
```

```
## iter 120 value 6603.485329
```

```
## iter 130 value 6603.453476
```

```
## iter 140 value 6603.450232
```

```
## final value 6603.449303
```

```
## converged
```

```
train$pred_nn10 <- predict(nn_fit10, type = "r")
```

```
auc(train$train.y, train$pred_nn10)
```

```
## Area under the curve: 0.79
```

The neural network model significantly improves the AUC from the model with the logistic regression to the ANN with two hidden layers to the ANN with five hidden layers. There is an improvement from the

nnet model with five hidden layers to the model with ten hidden layers as well. It is interesting that there is not an improvement from the logistic model to the nnet models with one layer.

- (d) Is there evidence for overfitting in your results in (c)? Use the test data, also available in `mnist_all.RData`, to find out.

```
test$pred_nn2 <- predict(nn_fit2, newdata = test, type = "r")
```

```
auc(test$test.y, test$pred_nn2)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated
## use a matrix as predictor. Unexpected results may be produced, please pass
## a numeric vector.
```

```
## Area under the curve: 0.787
```

```
test$pred_nn5 <- predict(nn_fit5, newdata = test, type = "r")
```

```
auc(test$test.y, test$pred_nn5)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated
## use a matrix as predictor. Unexpected results may be produced, please pass
## a numeric vector.
```

```
## Area under the curve: 0.791
```

```
test$pred_nn10 <- predict(nn_fit10, newdata = test, type = "r")
```

```
auc(test$test.y, test$pred_nn10)
```

```
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated
## use a matrix as predictor. Unexpected results may be produced, please pass
## a numeric vector.
```

```
## Area under the curve: 0.793
```

The auc for the test set is not drastically different from that of the train set. For the hidden node with two layers, the test set does only slightly worse. The auc for the test set actually does better for five nodes and ten nodes. There is therefore not significant evidence for overfitting. As indicated by the AUCs on the train sets, the bias was not that low, so we should not expect that high of a variance.

5 points each:

Extra 36 (5 points)

In the Tensorflow Playground, we can use a “bullseye” dataset to demonstrate non-linear decision boundaries that would be impossible for logistic regression. Here we’re going to explore that kind of dataset in a simplified version. A one-dimensional bullseye dataset would be like the following. Notice that one of the classes completely surrounds the other.


```
x = rnorm(50, 0, 2)
y <- rep(1, length(x))
y[abs(x) < 1] = 0

plot(x, rep(0, length(x)), col = y + 1)
```


Figure 3: Alt text

(a) Fit a logistic regression model to this dataset. Verify that the results are not great.

```
set.seed(1)
x = rnorm(50, 0, 2)
y<-rep(1, length(x))
y[abs(x) < 1] = 0
plot(x,rep(0,length(x)),col=y+1)
```



```
data <- data.frame(x, y)
fit <- glm(y ~ x, data = data, family=binomial)
data$pred <- predict(fit, type = 'response')
r <- roc(data$y, data$pred)
plot(r)
```




```
auc(data$y, data$pred)
```

```
## Area under the curve: 0.643
```

The AUC of about .643 is close to the point at .50 where just switching all of our classifications would improve the classification. The ROC curve also shows clearly that the results are not great.

- (b) But we can solve this problem using logistic regression if we employ clever “feature engineering”. Create a new feature which is just x^2 . Make a plot of the two features x and x^2 and color by class label to verify that the two classes are now more easily separable. Fit a logistic regression model and comment on the results.

```
data$x2 <- x^2
ggplot(data, aes(x = x, y = x2, color = factor(y))) + geom_point(size = 1, alpha = 1, na.rm = TRUE)
```




```
fit <- glm(y ~ x + x2, data = data, family=binomial)
summary(fit)
```

```
##
## Call:
## glm(formula = y ~ x + x2, family = binomial, data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.37e-04 -2.00e-08  2.00e-08  2.00e-08  1.34e-04
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -147.3     39061.7      0       1
## x              13.7     25374.6      0       1
## x2            121.4     30656.2      0       1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6.8593e+01  on 49  degrees of freedom
## Residual deviance: 3.6937e-08  on 47  degrees of freedom
## AIC: 6
##
## Number of Fisher Scoring iterations: 25
```

```
#
data$pred_2 <- predict(fit, type = 'response')
```

```
#
r <- roc(data$y, data$pred_2)
plot(r)
```



```
auc(data$y, data$pred_2)
```

```
## Area under the curve: 1
```

According to the ROC and AUC, this model predicts classifications perfectly. Some observations ended up with predicted probabilities of 1. It is important to note that neither X nor X2 are significant. In fact, they are highly insignificant, meaning that we've ended up with very high variability and very low bias (i.e. overfitting).

- (c) If we never thought of this feature engineering, we can also easily solve this problem with a neural network. But importantly, we have to make a network topology such that the hidden layer has higher dimensionality than the input layer. Fit a neural network to Y ??? X with two nodes in the hidden layer. Verify that we can achieve perfect classification on the training data.

```
nn_fit <- nnet(y ~ x, data= data, maxit = 2000, decay = .01, size=2)
```

```
## # weights: 7
## initial value 13.398152
## iter 10 value 8.893467
## iter 20 value 8.637542
## iter 30 value 5.236470
## iter 40 value 3.408634
## iter 50 value 3.357195
## final value 3.357193
```

```
## converged
data$pred_nn <- predict(nn_fit, type = "r")
auc(data$y, data$pred_nn)
```

Area under the curve: 1

We achieve perfect classification on the data through this method.

- (d) By projecting the data into a higher-dimensional space, we can separate the two classes. In the case of a neural network, the network figured it out for us - we didn't have to do it ourselves. Provide an explanation and intuition into how the network can achieve this goal in this particular case. Your explanation might rely on helpful visualizations.

When the two classification classes are linearly separable, then we can just use a logistic regression, which uses hyperplanes as decision boundaries. However, when we have nonlinear decision boundaries like we do in this case, neural networks are a useful tool because it uses matrix multiplication to identify a matrix of weights that minimizes the squared loss, as opposed to maximizing likelihood, through nonlinear functions. Any continuous function of n variables can be approximated arbitrarily well by a feed-forward artificial neural network with one hidden layer and finitely many layers (two nodes in that layer in this case) according to the universal approximation property of neural networks. Our model is therefore

Extra 37 (5 points)

The data for this exercise are in the UCI Machine Learning Repository, <http://archive.ics.uci.edu/ml/index.php>. We shall use the concrete compressive strength data. For details and a link, refer to problem 11.

- a) Import the data into your R workspace and change all variable names to something simpler. Split the data into a training set (70%) and a test set (30%).

```
concrete <- read_excel("Concrete_Data.xls")
colnames(concrete) <- c("cementkg", "blustfur", "flyash", "superplas", "courseagg", "fineagg", "age", "compressive_strength")
# set the seed to make your partition reproducible
set.seed(246)
smp_size <- floor(0.80 * nrow(concrete))
train <- sample(seq_len(nrow(concrete)), size = smp_size)
concrete_train <- concrete[train,]
concrete_test <- concrete[-train,]
```

- b) Fit artificial neural networks with a single hidden layer and 2, 3, 4, . . . 20 nodes to the training data. Compute the root mean squared residuals for each network and plot this quantity against the number of hidden nodes.

```
train_rmsr <- numeric()
for(i in 1:20) {
  nn_fit <- nnet(CCS ~ ., data= concrete_train, maxit = 2000, decay = .01, size = i)
  train_rmsr[i] <- sqrt(mean((predict(nn_fit, type = "r") - concrete_train$CCS)^2))
}
```

```
## # weights: 11
## initial value 4719605.846369
## iter 10 value 4693761.066038
## final value 4693759.090905
## converged
```

```

## # weights: 21
## initial value 4720077.112824
## iter 10 value 4693761.572103
## final value 4693759.038919
## converged
## # weights: 31
## initial value 4740912.271542
## iter 10 value 4693789.989285
## iter 20 value 4693760.163129
## final value 4693759.305674
## converged
## # weights: 41
## initial value 4727684.839965
## iter 10 value 4694207.975873
## iter 20 value 4693774.087426
## iter 30 value 4693759.097906
## final value 4693758.936661
## converged
## # weights: 51
## initial value 4747977.363951
## iter 10 value 4693784.931301
## iter 20 value 4693759.213916
## final value 4693758.813183
## converged
## # weights: 61
## initial value 4703543.097364
## iter 10 value 4693759.666783
## final value 4693758.569362
## converged
## # weights: 71
## initial value 4730954.905412
## iter 10 value 4693759.506057
## iter 10 value 4693759.479089
## final value 4693759.035354
## converged
## # weights: 81
## initial value 4730742.229266
## iter 10 value 4693761.039890
## final value 4693758.749623
## converged
## # weights: 91
## initial value 4744621.415970
## iter 10 value 4694964.487179
## iter 20 value 4693806.522970
## iter 30 value 4693759.188802
## iter 30 value 4693759.148989
## iter 30 value 4693759.102746
## final value 4693759.102746
## converged
## # weights: 101
## initial value 4750616.115699
## iter 10 value 4693766.664089
## iter 20 value 4693758.931975
## final value 4693758.377477

```

```


## converged
## # weights: 111
## initial value 4716308.812495
## iter 10 value 4693931.692275
## final value 4693758.655269
## converged
## # weights: 121
## initial value 4715390.309844
## iter 10 value 4693765.900488
## final value 4693759.458032
## converged
## # weights: 131
## initial value 4712503.135523
## iter 10 value 4693767.928268
## final value 4693758.395214
## converged
## # weights: 141
## initial value 4744615.934732
## iter 10 value 4695365.125791
## iter 20 value 4693795.850849
## iter 30 value 4693758.448772
## final value 4693758.253790
## converged
## # weights: 151
## initial value 4705696.826160
## iter 10 value 4696051.134312
## iter 20 value 4694107.576509
## final value 4693758.352035
## converged
## # weights: 161
## initial value 4760700.179369
## iter 10 value 4693775.189426
## iter 20 value 4693758.513553
## final value 4693758.320165
## converged
## # weights: 171
## initial value 4743111.257917
## iter 10 value 4695257.462941
## iter 20 value 4693810.594291
## final value 4693758.928702
## converged
## # weights: 181
## initial value 4724052.983660
## iter 10 value 4693759.722584
## final value 4693758.329378
## converged
## # weights: 191
## initial value 4736266.299586
## iter 10 value 4695550.374578
## iter 20 value 4693775.252083
## iter 30 value 4693760.421962
## final value 4693758.371889
## converged
## # weights: 201

```



```
## initial value 4700409.238789
## iter 10 value 4693767.726971
## iter 20 value 4693758.273062
## iter 20 value 4693758.259355
## iter 20 value 4693758.248455
## final value 4693758.248455
## converged

train_rmsr <- data.frame(cbind(1:20, train_rmsr))
ggplot(train_rmsr, aes(x = V1, y = train_rmsr)) + geom_point(size = 1, alpha = 1, na.rm = TRUE)
```



c) For the networks in b), compute also the root mean squared residuals on the test data and plot them in the same graph.

```
test_rmsr <- numeric()

for(i in 1:20) {
  nn_fit <- nnet(CCS ~ ., data= concrete_train, maxit = 2000, decay = .01, size = i)
  test_rmsr[i] <- sqrt(mean((predict(nn_fit, newdata = concrete_test, type = "r") - concrete_test$CCS)^2))
}
```

```
## # weights: 11
## initial value 4729302.365435
## iter 10 value 4694227.216250
## iter 20 value 4693775.670164
## final value 4693760.472159
## converged
## # weights: 21
```

```

## initial value 4710799.743444
## iter 10 value 4694012.589964
## iter 20 value 4693773.543211
## iter 30 value 4693759.153074
## final value 4693759.055387
## converged
## # weights: 31
## initial value 4719422.126321
## iter 10 value 4694359.610637
## iter 20 value 4693784.513689
## final value 4693759.009337
## converged
## # weights: 41
## initial value 4730838.963614
## iter 10 value 4693835.444067
## iter 20 value 4693758.855832
## final value 4693758.746229
## converged
## # weights: 51
## initial value 4740178.195552
## iter 10 value 4694735.085502
## iter 20 value 4693804.246611
## iter 30 value 4693760.838899
## final value 4693759.306689
## converged
## # weights: 61
## initial value 4730198.985032
## iter 10 value 4694725.574800
## iter 20 value 4693803.296897
## iter 30 value 4693759.157988
## final value 4693758.715067
## converged
## # weights: 71
## initial value 4723961.866460
## iter 10 value 4693764.683026
## iter 20 value 4693759.108100
## final value 4693758.717296
## converged
## # weights: 81
## initial value 4735191.344958
## iter 10 value 4693759.338280
## final value 4693758.408167
## converged
## # weights: 91
## initial value 4719882.126359
## iter 10 value 4693767.940595
## final value 4693758.496051
## converged
## # weights: 101
## initial value 4744872.587199
## iter 10 value 4694906.345843
## iter 20 value 4693831.173164
## iter 30 value 4693759.391376
## final value 4693758.544459

```

```


## converged
## # weights: 111
## initial value 4743757.813350
## iter 10 value 4693764.333856
## final value 4693758.455412
## converged
## # weights: 121
## initial value 4724116.759807
## iter 10 value 4693758.876623
## iter 10 value 4693758.856149
## iter 10 value 4693758.821910
## final value 4693758.821910
## converged
## # weights: 131
## initial value 4737240.418721
## iter 10 value 4693873.534684
## iter 20 value 4693762.516431
## final value 4693758.454293
## converged
## # weights: 141
## initial value 4720331.678715
## iter 10 value 4693769.375464
## final value 4693759.262788
## converged
## # weights: 151
## initial value 4714594.401830
## iter 10 value 4695468.759959
## iter 20 value 4693782.548144
## final value 4693758.513508
## converged
## # weights: 161
## initial value 4739409.487775
## iter 10 value 4693760.490039
## final value 4693758.552958
## converged
## # weights: 171
## initial value 4709730.224177
## iter 10 value 4693833.117051
## final value 4693758.540616
## converged
## # weights: 181
## initial value 4726750.545799
## iter 10 value 4693761.121054
## final value 4693758.283411
## converged
## # weights: 191
## initial value 4717060.392373
## iter 10 value 4693763.862299
## final value 4693758.360945
## converged
## # weights: 201
## initial value 4711377.296717
## iter 10 value 4693811.257787
## final value 4693758.211858

```

```
## converged
rmsr <- data.frame(cbind(train_rmsr, test_rmsr))

# melt the data to a long format
rmsr <- melt(data = rmsr, id.vars = "V1")

# plot
ggplot(data = rmsr, aes(x = V1, y = value, colour = variable)) + geom_point()
```



d) Is there evidence of overfitting? How can you tell?

When placing the RMSR for the train and test on the same plot, it is clear that there is an overfitting problem. The test set RMSR is consistently higher than the train RMSR, meaning that our model performs worse when applied to new data.

e) Do you think that the ANN is overfitting the data?

ANN appears to be overfitting the data. Even when there is just one node, the RMSR is much higher than that of the test set.